

Anix Lynch

Culver City, CA • anix@anixlynch.com • (424) 410-0672
gozeroshot.dev • linkedin.com/in/anixlynch • github.com/anix-lynch
Authorized to work in the US (Green Card)

PROFESSIONAL SUMMARY

Hands-on AI & Agentic Systems Engineer and former Venture Capital General Partner combining deep investment-management expertise with the ability to architect, build, evaluate, and deploy AI applications end to end. Brings experience deploying \$15M in venture capital, evaluating 700+ technology companies annually, and supporting institutional investment operations across platforms with approximately \$2B in AUM. Builds production-style AI systems using Python, FastAPI, LangGraph, RAG, MCP, and cloud infrastructure, with measurable results in decision accuracy, retrieval quality, operational resilience, and governed analytics. Uniquely positioned to translate complex business requirements into auditable, decision-ready AI solutions—bridging executive judgment, financial domain expertise, and hands-on engineering execution.

EDUCATION

UNIVERSITY OF CHICAGO BOOTH SCHOOL OF BUSINESS

MBA

Chicago, IL

RITSUMEIKAN UNIVERSITY

BA, International Business

Kyoto, Japan

EXPERIENCE

ZEROSHOT (INDEPENDENT CONSULTANT)

AI / Agentic Systems Engineer

02/2023 – Present

- **AI Research & RAG — Evidence Reliability:** Built a citation-grounded AI research service using Python, FastAPI, RAG, Gemini, and retrieval evaluation, achieving Hit@5 of 0.95, MRR of 0.90, and NDCG@10 of 0.89—enabling traceable, evidence-backed analysis and reducing reliance on unsupported AI-generated answers. Metrics were reproduced across a 20-query BM25 evaluation; 12 targeted API and retrieval tests passed.
- **Agentic AI — Controlled Decision Automation:** Engineered LangGraph-based agentic workflows with tool calling, human-in-the-loop (HITL) approvals, bounded retries, and idempotent execution—achieving 94% decision correctness and 90% incident recovery in evaluated scenarios to support reliable automation of complex, high-stakes business processes. Separate suites measured 47/50 labelled decisions and recovery of 63/70 recoverable simulated incidents; 14 action-engine tests passed.
- **MCP & Enterprise Integration — Operational Scale:** Implemented FastMCP servers, tool-calling interfaces, and API integrations, validating end-to-end tool discovery and execution; separately developed AI-assisted Airtable workflows covering 1,354 contacts across 525 companies, demonstrating reusable integration patterns for scalable research and operational coordination.
- **AI Quality & Security — Governed Automation:** Built LLM evaluation and guardrail controls covering prompt injection, PII masking, unsafe actions, groundedness, and failure handling, supported by automated acceptance tests and human approval gates—enabling safer, auditable AI-assisted workflows with controlled access to sensitive information. Blocked 6/6 known-pattern injection tests, with broader paraphrase gaps explicitly tracked in the red-team suite.
- **Microsoft Fabric & Power BI — Decision Intelligence:** Built governed analytics pipelines using Microsoft Fabric, OneLake, Direct Lake, Power BI, DAX, and related TMDL source, reconciling four table counts across GCP and Fabric and achieving sub-1.3-second DAX execution—delivering consistent, responsive reporting for data-driven business decisions. Across ten benchmark runs, p50 latency was 0.54s and 0.56s and maximum latency was 0.91s.
- **Cloud Deployment & CI/CD — Operational Reliability:** Deployed a live, containerized Python/FastAPI AI application using Docker, Google Cloud Run, and CI/CD, incorporating health checks, automated testing, and release controls—demonstrating repeatable deployment and operational readiness beyond notebook prototypes. The public deployment is live, and 12 targeted API and retrieval tests passed.
- **Multi-Model AI Engineering — Flexible Enterprise Architecture:** Integrated Gemini/Vertex AI into grounded retrieval and evaluation workflows, with a private Claude/Anthropic implementation, Bedrock Titan implementation, and OpenAI configuration evidence documented across projects; combined API-based model access with testable orchestration patterns to support adaptable AI applications without dependence on a single model provider, without claiming a single three-provider production router.

EXPARA VENTURE (SINGAPORE-BASED SEA VC)

General Partner

05/2017 – 05/2020

- Screened 700+ AI, DeepTech, FinTech, and infrastructure startups per year and governed KPI reporting for 7 portfolio companies, translating technical, operating, and financial signals into investment and board decisions.
- Structured and launched Thailand's first institutional, government-backed seed VC matching fund; deployed \$15M in anchor capital and syndicated private co-investors.
- Served as board director for 7 portfolio companies, driving milestones, runway oversight, operational problem solving, and follow-on fundraising.

MQDC, CP GROUP FAMILY OFFICE

VP Business Intelligence / Chief of Staff / Innovation Strategy Lead

12/2012 – 06/2016

- Built BI and data-pipeline governance for a national innovation strategy, consolidating fragmented data across 26 programs into executive decision systems.
- Modeled approximately \$395M in capital requirements; 17 of 26 initiatives entered implementation within three years, with policy alignment contributing to \$560M+ in public innovation funding.
- Originated the True Digital Park concept and structured investment theses across smart cities, alternative protein, vertical farming, and a \$162M aging-tech feasibility program.

CAPITAL REALTY LLC (ANA CROWNE PLAZA KOBE)

Asset Manager

08/2009 – 02/2011

- Led a \$400–500M renovation program and consolidated tenant sales, occupancy, and foot-traffic data across 75+ property managers.
- Built transaction-driven tenant optimization models over 10K+ monthly records, lifting retail revenue approximately 15% and reducing budget overruns from 12% to 3%.

MGPA (BLACKROCK REAL ESTATE PLATFORM)

Investment Associate

04/2008 – 06/2009

- Reconciled multi-source performance datasets to identify 15–25% underwriting risk variances across approximately \$2B AUM funds.
- Built structured underwriting tools that reduced modeling errors approximately 25% and accelerated diligence from three weeks to ten days.

NEW CITY CORPORATION

BizOps / CEO Office

10/2006 – 04/2008

- Led production of 50+ PPMs and institutional pitch books within an \$800M–\$1.6B AUM platform, consolidating multi-asset datasets for investor reporting.

APPLE INTERNATIONAL COMPANY LTD

BizOps / Strategic Planning

04/2001 – 09/2006

- Led cross-border financial consolidation across 5 APAC entities with an approximately \$250–350M revenue base, supporting IPO execution and institutional reporting.
- Built reconciliation frameworks that held metric integrity above 99% and reduced reporting latency approximately 30%.

SKILLS

Generative AI & Agents: RAG, LangChain, LangGraph, tool calling, multi-step orchestration, persistent state, human-in-the-loop, MCP

Evaluation & Safety: groundedness and retrieval evaluation, Hit@K / MRR / NDCG, prompt-injection tests, PII masking, citation validation, release gates, replayable receipts

Engineering & Cloud: Python, FastAPI, SQL, REST APIs, Docker, Terraform, GitHub Actions, GCP / Vertex AI, AWS Bedrock, Microsoft Azure

Data & Microsoft: Microsoft Fabric, OneLake, Power BI Direct Lake, DAX, BigQuery, dbt, data contracts, semantic layers, data governance

CERTIFICATIONS

- DeepLearning.AI Data Engineering Professional Certificate (2026) — 10 AWS services, 63 Terraform modules, Great Expectations, Kinesis streaming
- IBM AI Developer Specialization (2025) — RAG, Agentic AI, LangChain, Vector Databases
- IBM Generative AI Engineering (2025) — Transformers, NLP, Prompt Engineering
- Deep Learning Specialization (Stanford/Coursera) — Neural Networks, PyTorch, TensorFlow